

THE CHOICE LINE

AI is present. People retain the choice.

Others test whether AI can be withdrawn. Choice Line tests whether a person can finish the same task without being penalized for declining AI.

1 — 3 — 2 — 12 — 6 — 1

No penalty for declining AI

Service Twin turns choice into a deliverable.

1

Choice first

See both routes before collection

2

Same outcome

One task and one accepted outcome

3

Completable

Offline, staffed, and phone-free completion

4

Takeover

Immediate human takeover on failure

5

Appealable

Receipt, owner, and response time

6

Measured


Proposed scale gate: CEI ≥ 90 and each item ≥ 80

One evidence chain, three scales

One evidence chain, three scales

From public task to reviewable space, without disguising rough geometry as a precise basemap

Overview map • provisional scope



Three scopes and evidence chain

43.6 km²

Research • announcement text

≈11.4 km²

Overall design • announced value

368.4 ha

Three key areas • announced sum

- 1 Announcement / taskbook
- 2 Repository provisional geometry
- 3 Nine GeoJSON layers
- 4 EPSG:4548 metrics
- 5 Scenarios and risks
- 6 Bilingual figures / PDF

All alignments await survey and rights checks

Source: submission GeoJSON, metrics.json, source/assumption and scenario models; EPSG:4548 recomputation

Provisional boundary for proposal generation and self-check only; not a redline, survey, title, approval, or engineering basis

The 1-3-2-12-6-1 framework


02/10

THE CHOICE LINE

The 1-3-2-12-6-1 spatial armature

One walk links three labs, two wings, twelve nodes, six safeguards, and one disclosure system

Land-use partition • topologically closed



Source: submission GeoJSON, metrics.json, source/assumption and scenario models; EPSG:4548 recomputation

Spatial mnemonic

1-3-2-12-6-1

1

Ground Truth Walk

3

choice labs

2

collaborative wings

12

Service Twin nodes

6

choice safeguards

1

CEI disclosure system

Jing-Zhang culture

urban AI life

AI-integrated innovation

Provisional boundary for proposal generation and self-check only; not a redline, survey, title, approval, or engineering basis

02

Three zones form an evidence pipeline

03/10

THE CHOICE LINE

Three zones form one evidence pipeline

The north discovers failure, the middle enables public review, and the south accepts real services

NORTH • ZHONGZHIYUAN

FAIL-SAFE GARDEN

- 1 Robot safe stop and human handover**
Trial operator and on-site safety marshal
- 2 Model red-team and failure drill**
Independent assurance lead and service owner
- 3 Camera-free edge perception benchmark**
Data steward and facilities operator
- 4 Enterprise Agent interoperability and handover**
Enterprise service owner, interface engineer, and human handover...

Output: reproducible failure + takeover

PROVISIONAL RECTANGLE

MIDDLE • AI ORIGIN

OPEN LEARNING COMMONS

- 5 Account-free learning and research explanation**
Public-learning curator and duty explainer
- 6 Health-service navigation without diagnosis**
Health-navigation owner and qualified service worker
- 7 Legal and public-rule explanation without adjudication**
Public-service owner, rule editor, and independent reviewer
- 8 Accessible service co-test**
Accessibility lead, community representative, and facilities owner

Output: understandable rules + public review

PROVISIONAL RECTANGLE

SOUTH • DAZHONGSI

CIVIC CHOICE HUB

- 9 Station-park transfer navigation**
Mobility-service owner and on-site transfer steward
- 10 Bilingual Jing-Zhang cultural guide**
Cultural curator, history editor, and on-site interpreter
- 11 Anonymous and non-personalized retail**
Retail operator, consumer-rights contact, and shop worker
- 12 Face-free event flow**
Event owner, facilities maintainer, data steward, and safety marshal

Output: real tasks + route gap

PROVISIONAL RECTANGLE

Discover failure → public review → real-service acceptance → anonymized feedback

Source: submission GeoJSON, metrics.json, source/assumption and scenario models; EPSG:4548 recomputation

Provisional boundary for proposal generation and self-check only; not a redline, survey, title, approval, or engineering basis

One walk, equally visible twin interfaces


04/10

THE CHOICE LINE

One walk, two equally visible interfaces

Seven ports and three sections join choice, access, blue-green space, and service continuity

Slow mobility + blue-green public space

**9.72 km**

概念步道 / conceptual walk

7

缝合门户 / stitching ports

15.29%

提案绿脊 / proposal green spine

Three standard sections

Sensed / camera-free



notice before sensing + equal bypass

Smart / offline



screen, print, and tactile at equal prominence

Automatic / human



one case number, same price and hours

The task remains completable during network loss

Source: submission GeoJSON, metrics.json, source/assumption and scenario models; EPSG:4548 recomputation

Provisional boundary for proposal generation and self-check only; not a redline, survey, title, approval, or engineering basis


04

Choice equivalence must be measured

Choice equivalence must be measured

Six CEI dimensions, geometry recalculation, and compliance evidence govern expansion

CEI • Choice Equivalence Index



proposal target • pending baseline

Core metrics + evidence closure

11.413 km²
provisional site area

15.29%
proposal green ratio

0.89%
node public space

12 / 7 / 4
nodes / ports / landmarks



unknown is not zero: statutory FAR remains unknown

- Any component < 80: do not scale
- Weakest persona fails: no average override
- No funded human shift: no Service Twin
- Official polygon arrives: recalculate the package

Source: submission GeoJSON, metrics.json, source/assumption and scenario models; EPSG:4548 recomputation

Provisional boundary for proposal generation and self-check only; not a redline, survey, title, approval, or engineering basis


One task, two completable routes

06/10

THE CHOICE LINE

One task, two completable routes, one owner

Service Twin is a pilot deliverable, not a hidden fallback button



Four-party loop: AI innovation does not automate accountability



Source: submission GeoJSON, metrics.json, source/assumption and scenario models; EPSG:4548 recomputation

Provisional boundary for proposal generation and self-check only; not a redline, survey, title, approval, or engineering basis

06

Twelve scenarios that can stop

07/10

THE CHOICE LINE

Twelve scenario cards, twelve stoppable commitments

Every card fixes a data budget, human takeover, failure, and appeal

S01Robot safe stop and human handover

DATAProcess sensors on device without uploading raw imagery; propose a 30-day minimal incident log subject to...

TWINA human courier completes the same task on the same accessible route with fee and hours displayed together

STOPStop immediately and isolate equipment after boundary breach, loss of link, or unrecognized works or obstacle

INDUSTRY ASSURANCErobot-handover

S02Model red-team and failure drill

DATAUse no real personal data; de-identify logs and delete them under the project test plan

TWINIndependent experts apply the same public checklist and publish findings beside model-based assurance

STOPPause after test-set contamination, irreproducible results, or escape of high-risk output

INDUSTRY ASSURANCEmodel-red-team

S03Camera-free edge perception benchmark

DATACollect no face, voiceprint, or device identifier and publish aggregate counts only

TWINTrained observers record the same non-identifying statistics during matched periods

STOPStop publication after identity inference, unsafe blind spots, or undersized aggregate samples

INDUSTRY ASSURANCEcamera-free-perception

S04Enterprise Agent interoperability and handover

DATAUse synthetic or cleared data only; tokens are short-lived and immediately revocable

TWINA human specialist completes the task using the same case number and exported material

STOPStop the pilot when export omits fields, permissions cannot revoke, or a person cannot restore service

INDUSTRY ASSURANCEagent-interoperability

S05Account-free learning and research explanation

DATAErase on-device preference at session end and build no learning profile

TWINA staffed explanation, print reader, and device-free labels deliver the same learning goal

STOPWithdraw content after missing sources, claims beyond evidence, or absent accessible format

PUBLIC SERVICEaccount-free-learning

S06Health-service navigation without diagnosis

DATAStore no symptoms and infer no disease; send appointment essentials directly to the accountable...

TWINPhone, print directory, and trained staff complete the same search and appointment help

STOPTransfer immediately after diagnostic language, stale provider data, or urgent risk

PUBLIC SERVICEhealth-navigation

S07Legal and public-rule explanation without adjudication

DATADo not request unnecessary identifiers and erase temporary context when the session ends

TWINAn in-person desk, phone line, and print checklist explain the same case under one identifier

STOPStop automated response after uncertain version, conflicting rules, or a rights determination

PUBLIC SERVICERULES-explainer

S08Accessible service co-test

DATACollect no biometric data; de-identify feedback and permit withdrawal

TWINPhysical wayfinding, staffed accompaniment, and a print checklist complete the same task

STOPPause after a critical route break, non-withdrawable feedback, or inconsistent safety information

PUBLIC SERVICEaccessible-co-test

S09Station-park transfer navigation

DATANo persistent location tracking; erase on-device location at session end

TWINPhysical signs, a paper route map, and staffed help complete the same arrival task

STOPStop recommendations after lift failure, stale information, or an inaccessible route break

PUBLIC SERVICEstation-park-transfer

S10Bilingual Jing-Zhang cultural guide

DATABuild no visitor profile and erase optional session preferences at the end

TWINPhysical labels, paper map, and a human guide independently deliver the full narrative

STOPWithdraw after unknown provenance, distorted translation, or describing a proposal as implemented

PUBLIC SERVICEjingzhang-cultural-guide

S11Anonymous and non-personalized retail

DATACollect no face, device identifier, or cross-store history and erase the session at exit

TWINA staff member and paper category index complete shopping at the same price and hours

STOPDisable after price disparity, sensitive inference, manipulative interface, or failed session deletion

PUBLIC SERVICEanonymous-retail

S12Face-free event flow

DATARetain time-banded aggregate counts only and delete raw counts after the event under the published rule

TWINHuman patrol, physical zone signs, and public address deliver the same safety operation

STOPDisable after distorted counts, storm threshold, blocked accessible exit, or failed deletion

PUBLIC SERVICEface-free-event-flow

Source: submission GeoJSON, metrics.json, source/assumption and scenario models; EPSG:4548 recomputation

Provisional boundary for proposal generation and self-check only; not a redline, survey, title, approval, or engineering basis

The least-served persona gates expansion

08/10

THE CHOICE LINE

The weakest persona governs scale

An average cannot hide failure among older, blind, minor, or international users

Eight personas are acceptance tasks, not labels

1

P01 • Older adult or feature-phone user

Complete the service without an app, account, or QR code using large-print, phone, and in-person support.

Exclusion risk: Forced device purchase, longer queues, or denial because a digital credential is missing.

2

P02 • Commuter with visual, hearing, or mobility disability

A continuous accessible route, tactile and audio information, captions, seating, and a camera-free option.

Exclusion risk: Broken station-park continuity, hidden alternatives, or assistive and staffed services unavailable together.

3

P03 • Parent, caregiver, or minor

Shared plain-language notice before collection, clear guardian boundaries, care space, and no profiling of minors.

Exclusion risk: Default consent, manipulative interfaces, retained child data, or stroller routes omitted by automated advice.

4

P04 • Maintenance, cleaning, or delivery worker

Safe loading, rest and charging points, clear machine takeover rights, isolation procedures, and humane schedules.

Exclusion risk: Blocked by robot trials, uncompensated takeover labor, or penalties from performance models without appeal.

5

P05 • International visitor or developer

Bilingual rules, short-term account-free access, cross-cultural explanation, and clear participation and exit limits.

Exclusion risk: Chinese-only digital interfaces, inaccessible local accounts, or mistaking an exhibition for official approval.

6

P06 • University student or researcher

Open evaluation, quiet study, explainable research, night safety, exercise, and affordable daily services.

Exclusion risk: Innovation space reserved for corporate users, opaque results, or disconnected night routes and study facilities.

7

P07 • SME developer or startup team

Affordable interoperability tests, compliance help, independent assurance, prototyping, and a defined retirement interface.

Exclusion risk: Single-vendor lock-in, unaffordable repeated assurance, or stranded data and services after a pilot.

8

P08 • Human reviewer or public-service worker

Safe workstations, shift capacity limits, training, takeover authority, appeal logs, and protection from automation pressure.

Exclusion risk: Human fallback treated as limitless free labor while reviewers absorb model errors without stop authority or rest.

Weakest-persona gate

A ≠ average satisfaction

A compares the smaller and larger priority-persona success rates. If any group is punished by fallback, the node does not scale.



No component below 80

Source: submission GeoJSON, metrics.json, source/assumption and scenario models; EPSG:4548 recomputation

Provisional boundary for proposal generation and self-check only; not a redline, survey, title, approval, or engineering basis

Evidence first, anchors second, scale last

09/10

THE CHOICE LINE

Evidence first, then anchors, then scale

Three-stage gates, four-season operations, RACI, and human-route funding form one loop

Three stages: every expansion has a stop condition

P1 • 0–6 months

audit, baseline, bench tests
three no-permanent-work anchors

Gate: baseline + safety owner + OPEX

P2 • 6–18 months

one anchor per zone
continuous choice interface

Gate: CEI ≥ 90 and each component ≥ 80

P3 • 18–36 months

extend to twelve nodes
all else remains uncommitted

Gate: public review + safety + funding

Four-season operations

SPRING**GROUND TRUTH WEEK**

publish problems and baselines

SUMMER**FAIL-SAFE SEASON**

failure, storm, and handover

AUTUMN**OPEN SLEEPER FESTIVAL**

contribution and status review

WINTER**PUBLIC EVIDENCE ASSEMBLY**

scale, remedy, or retire

Responsibility and funding loop

- 1 service owner
- 2 site operator
- 3 independent evaluator
- 4 community
- 5 data steward
- 6 human reviewer

Free public base route ← cross-subsidized by
B2B assurance, training, and approved
events; no personal-data monetization

OPEX is an expansion gate

Source: submission GeoJSON, metrics.json, source/assumption and scenario models; EPSG:4548 recomputation

Provisional boundary for proposal generation and self-check only; not a redline, survey, title, approval, or engineering basis

09

From engineering autonomy to public autonomy


10/10

THE CHOICE LINE

From engineering autonomy to public autonomy

Open sleepers record contribution, version, license, and real adoption status—and permit correction

Centennial narrative: autonomy is not a one-off memorial



Others test whether AI can be withdrawn. Choice Line tests whether a person can finish the same task without being punished for declining AI.

不声称 ‘全球第一’ / NO WORLD-FIRST CLAIM

Four landmarks + open-sleeper state machine

LM-01 **Choice Mile Zero**
Extend the engineering autonomy of 1909 into public autonomy to understand, choose, and keep receiving service in the AI era.

LM-02 **Human Review Hall**
Turn model limits, appeals, takeover, and accountable ownership from a back-office process into visible civic space.

LM-03 **Fail-Safe Garden**
Publicly demonstrate when a system should stop, who can stop it, and how essential urban service continues afterward.

LM-04 **Open Sleeper Honor Line**
Record Agents, people, versions, licenses, and actual adoption status in an updateable and correctable form.

SUBMITTED → SELECTED → IMPLEMENTED

Correction can update credit, version, and state

Source: submission GeoJSON, metrics.json, source/assumption and scenario models; EPSG:4548 recomputation

Provisional boundary for proposal generation and self-check only; not a redline, survey, title, approval, or engineering basis

Eight global cases, no superficial copying

Each case separates fact, inference, and non-transferable boundary.

CASE 01

Amsterdam algorithm register and human appeal

Publish a decision card for every rights-affecting service and provide a human review and non-digital appeal route independent of the original automated result.

The mechanism relies on Dutch administrative law and identity infrastructure; registry disclosure is neither independent audit nor universal opt-out.

CASE 03

Singapore Global AI Assurance Sandbox

Use a deployer, site owner, and independent tester triad in Zhongzhiyuan, publishing risk hypotheses, failure thresholds, remediation, and retirement records.

The sandbox is neither a software runtime nor regulatory approval; technical testing does not replace process governance, legal review, or foundation-model safety.

CASE 05

MIT Kendall Square / InnovationHQ

Use a city-open ground floor and specialized upper-floor program stack in the Origin Community, funding operations alongside prototyping, IP, mentors, and capital services.

MIT's research density, brand, land, and capital network are not transferable, and the evidence does not show that buildings alone cause commercialization.

CASE 07

New York City statutory open-data governance

Assign a data steward for each participating district, campus, park, and operator, publishing classification, update cadence, field definitions, quality checks, and correction tickets.

Only legally publishable data is open; a portal cannot guarantee completeness or fitness and does not replace privacy review or sustained maintenance staff.

CASE 02

Norwegian data-protection regulatory sandbox

Create a compliance co-design room in the AI Origin Community where legal, technical, operational, and public representatives agree test plans and publish reusable exit reports.

Sandbox advice is not approval, certification, or a binding decision; participants retain compliance responsibility and the jurisdictional structure cannot be copied.

CASE 04

Estonian X-tee / X-Road exchange layer

Use source-retained data, service-specific authorization, revocable trust anchors, and verifiable call logs for cross-district cooperation rather than a single urban data lake.

X-Road is not identity, portal, or complete governance; endpoint authorization, data classification, and citizen-visible records still require separate design.

CASE 06

GOV.UK Assisted Digital and cross-channel service

Give phone, paper, in-person, and AI routes one service identifier; never hide the human channel, and feed offline failure evidence back into service improvement.

The guidance requires the channel mix users need, not every channel at all times, and disability access still needs its own research and formats.

CASE 08

Rotterdam Benthemplein water square

Design Dazhongsi event space for everyday use and stormwater storage, combining shade, seating, sediment cleaning, and programming in one maintenance brief.

Its storage capacity cannot be copied; runoff pollution, freeze-thaw, infiltration, child safety, and utility conditions require Beijing-specific calculation.

Twelve tasks, one delivery protocol

Data, twin route, stop rule, owner, and appeal are all mandatory.

S01

Robot safe stop and human handover

Complete low-speed delivery on a controlled park path and transfer safely to a person when risk triggers.

A human courier completes the same task on the same accessible route with fee and hours displayed together

S02

Model red-team and failure drill

Reproduce harmful output, leakage, bias, and failed takeover before deployment.

Independent experts apply the same public checklist and publish findings beside model-based assurance

S03

Camera-free edge perception benchmark

Test whether occupancy, obstacle, and queue assistance works without faces or identity.

Trained observers record the same non-identifying statistics during matched periods

S04

Enterprise Agent interoperability and handover

Prove a service can export, change provider, and transfer to a person after failure.

A human specialist completes the task using the same case number and exported material

S05

Account-free learning and research explanation

Help the public understand corridor AI work, evidence, and applicability limits.

A staffed explanation, print reader, and device-free labels deliver the same learning goal

S06

Health-service navigation without diagnosis

Explain public providers, opening hours, appointments, and accessible arrival options.

Phone, print directory, and trained staff complete the same search and appointment help

S07

Legal and public-rule explanation without adjudication

Turn public rules into understandable service steps and identify the accountable authority.

An in-person desk, phone line, and print checklist explain the same case under one identifier

S08

Accessible service co-test

Let disabled users co-accept routes, wayfinding, assistance, and switching between service modes.

Physical wayfinding, staffed accompaniment, and a print checklist complete the same task

S09

Station-park transfer navigation

Combine walking, cycling, and public transport without claiming an engineering alignment.

Physical signs, a paper route map, and staffed help complete the same arrival task

S10

Bilingual Jing-Zhang cultural guide

Connect the Jing-Zhang Railway, Zhan Tianyou, Zhongguancun innovation, and AI self-restraint.

Physical labels, paper map, and a human guide independently deliver the full narrative

S11

Anonymous and non-personalized retail

Compare product discovery without a personal profile against staffed shopping.

A staff member and paper category index complete shopping at the same price and hours

S12


Face-free event flow

Use aggregate counts to operate events, shade, and dual-mode stormwater space.

Human patrol, physical zone signs, and public address deliver the same safety operation

The human route has an owner and a budget

RACI, OPEX, and independent evaluation enter the scale gate.



Free public base route

Cross-subsidized by B2B testing, training, licensed events, and enterprise services; never by monetizing personal data.

Scale gate: baseline + safety owner + OPEX

Every claim can be traced backward

source → geometry → metric → scenario → figure → risk

11.413 km²

provisional site area

15.29%

proposal green ratio

0.89%

node public space

12 / 7 / 4

nodes / ports / landmarks



The claim ledger links every assertion to recomputable data and risk; unknown metrics remain unknown rather than becoming zero.

Unknown does not mean zero

Unknowns remain explicit for survey, clearance, baselining, and professional design.

UNKNOWN 1

Boundary and three key areas

Await registered official polygons

UNKNOWN 2

Approx. 9 km line

Requires measurement and definition clearance

UNKNOWN 3

Statutory FAR / height

Remain unknown; never infer

UNKNOWN 4

Hydraulics and pollution

Require specialist models and sampling

UNKNOWN 5

Human-route OPEX

Budget gate at every phase

UNKNOWN 6

CEI baseline

Target, not a current fact

What deserves memory must remain correctable

OPEN-SLEEPER STATE MACHINE



The monument records Agents, people, versions, licenses, and actual adoption state. Selection is not implementation; correction can update credit, version, and state.

NO WORLD-FIRST CLAIM