

Jing-Zhang Autonomy Commons: A Public Belt for the Autonomous-Mobility Era

Review delivery v4.0 · concept / provisional / unauthorised / not run / not a score

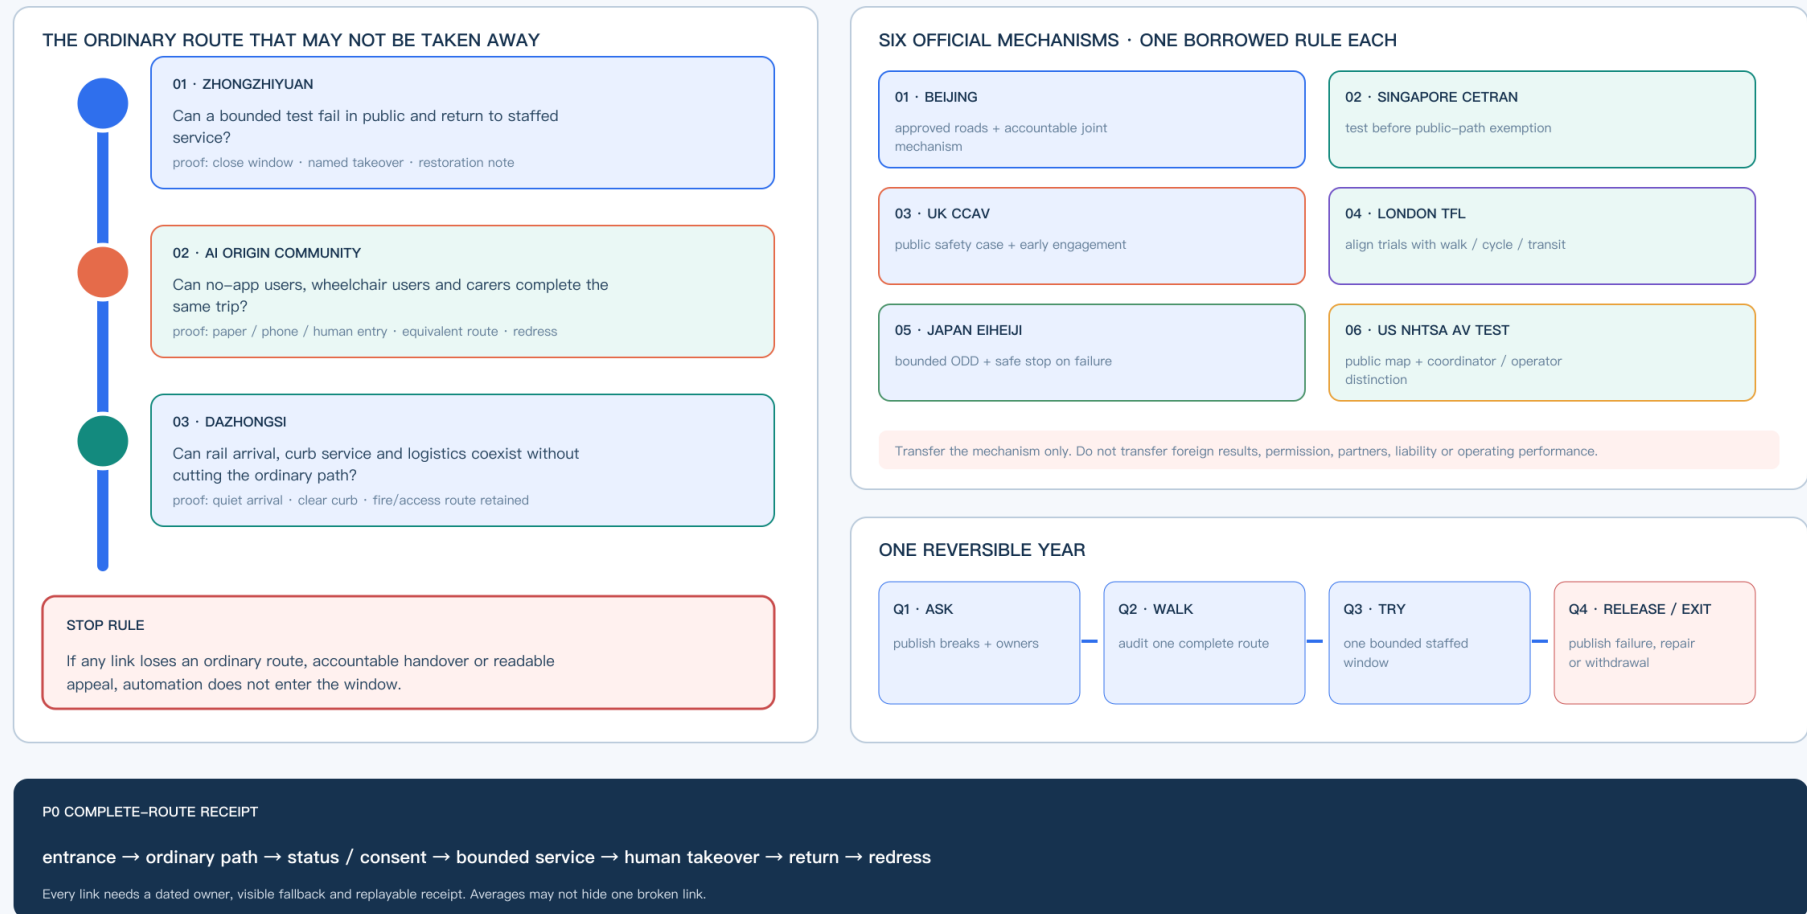
Autonomy Commons | prove the public route before admitting automation

One ordinary route · three public proof yards · six transferable mechanisms · one reversible year

AUTONOMY COMMONS · v4.0

CONCEPT / PROVISIONAL

01 / a spatial public-service argument, not a vehicle deployment plan



official_boundary=false · not_authorized_not_run · field_data=false · performance_results=null

mechanism transfer only · no foreign result imported

P0 COMPLETE-ROUTE RECEIPT: ONE BROKEN LINK FAILS THE TRIAL

entry → ordinary route → status / consent → bounded service → human takeover → return → redress

01 ORDINARY ROUTE complete even if the trial fails	02 HUMAN TAKEOVER named and staffed recovery	03 RETURN + REDRESS readable receipt; reopen
--	--	--

No approved boundary · no field data · no performance result; an average score cannot cancel a broken route.

Core proposition: the more automated driving spreads, the less the city should be designed around vehicles alone. The Jing-Zhang belt must protect continuous walking, wheelchair, cycling, child, care and maintenance routes before it decides where vehicles may operate.

This is an independent autonomous-mobility iteration, not a renamed copy of an earlier package. It uses the same provisional spatial base but has a different identity invariant: one ordinary-route continuity line, three test yards and two human-safety arcs. It adds machine-readable evidence for autonomous-driving curbs, low-speed shuttles, accessible service, remote intervention, data minimisation and failure rollback. Every spatial move remains a concept proposal; the package does not claim that any road inside the site is open to autonomous driving, nor that a vehicle, vendor or permit already exists [data:visual/assets/identity-system.json].

One-page executive brief: prove one public service chain before scaling automation

The first reversible acceptance unit is not “make the vehicle run.” It lets an ordinary person choose, use, stop, challenge and leave an automated service while returning to human service. This is a concept-level review interface only: it presupposes no measured length, road, vehicle or operator; exact points and engineering scale remain subject to formal boundaries, field walk-through and professional review.

Status: target design · not deployed · not authorized · not run

Caption: this board asks one public-space service question—before automation enters, can the ordinary route, human takeover, return journey and redress remain complete? The three nodes, six case mechanisms and Q1–Q4 loop are a concept review framework, not an approved road, deployed vehicle or field result.

Human-person path	Vehicle-space service	Evidence to leave	It tells
Choose ordinary or assisted route	Planned dual entry, target accessible surface, planned paper/phone entry	Choice state and obstruction record, without identity	No trial if ordinary route is unavailable
Request one low-speed service	Planned staffed desk, target legible curb, planned version status board	Scenario ID, version and equivalent human route	Hold if no equivalent human service exists
Trigger human takeover	Planned physical emergency stop, remote stop and local broadcast	Reason, time and responsible role	Stop automation if takeover is unavailable
Meet obstruction, network loss or weather	Planned isolation zone, rollback route and paper/phone fallback	Trigger, broadcast, evacuation and recovery record	Remain human-only; do not expand
Requester self-service replay and exit (not third-party verified)	Planned receipt, redress entrance and exit sign	Requester replay result, unproven items and closure record	Return to design stage if it cannot be replayed

The package currently has only an offline synthetic tabletop: four branches, seven contract checks, five rollback steps, and one synthetic negative replay for each stop_if. Each negative replay sends a trigger input through a pure decision path and records decision=reject_or_stop, result_status=not_run, and performance_results=null; it proves that the rejection/stop path is replayable at contract level, not field safety, authorization, deployment, or performance. not_authorized_not_run, performance_results=null, and the missing local baselines remain explicit [data:visual/assets/autonomy-curb-side-tabletop-contract.json#AUTONOMY-CURBSIDE-TABLETOP-001].

Identity system: make the right to exit visible

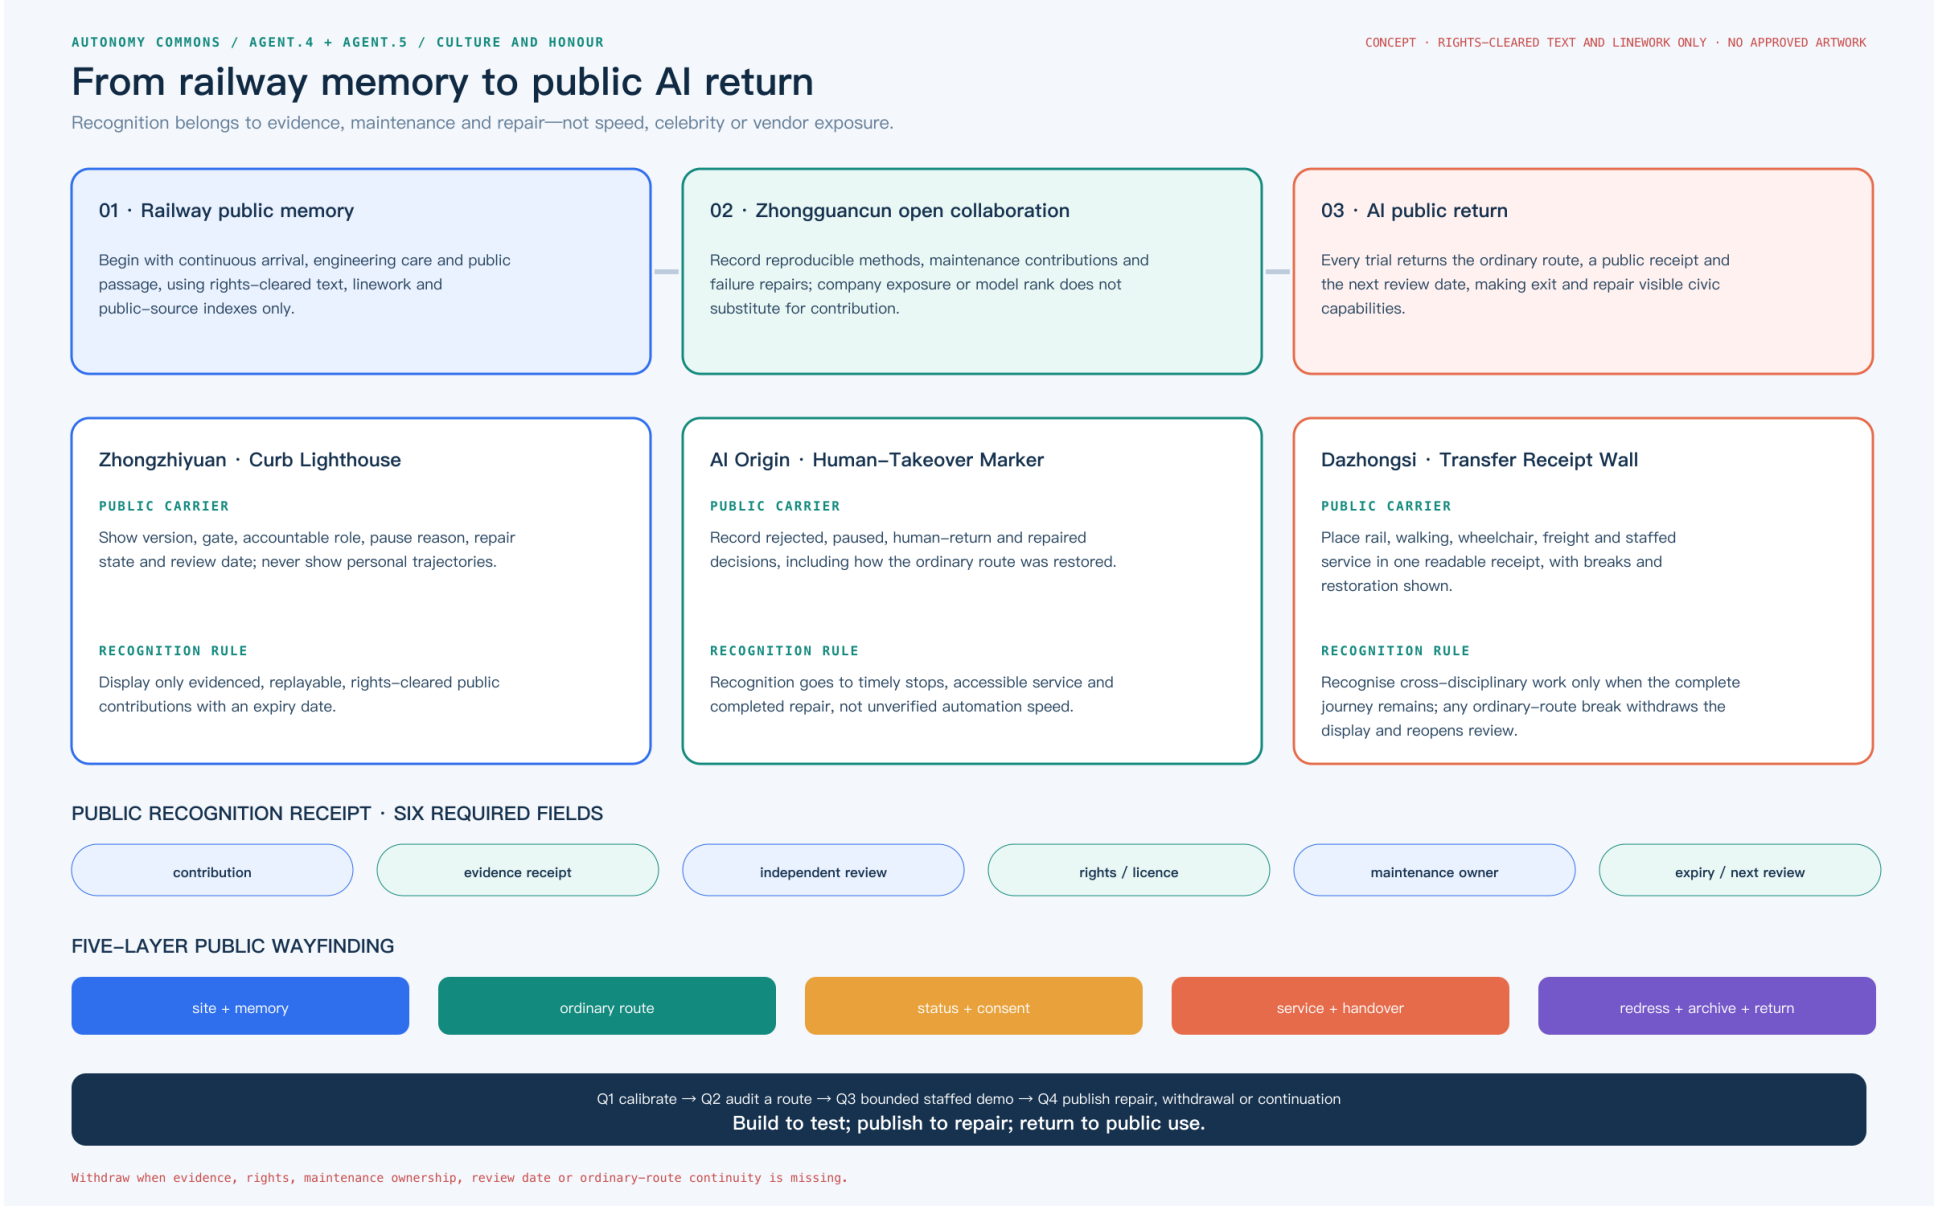
The “Jing-Zhang Autonomous Commons” mark is not a vehicle brand and does not reuse the previous parallel-rail identity. assets/identity/jingzhang-commons-mark.svg uses one ordinary-route line to connect three test yards and two human-safety arcs for takeover and reversible withdrawal; visual/assets/identity-system.json records high-contrast, tactile, paper, telephone and audible alternatives. The landmark catalogue grounds the identity in public interfaces: a Human-Takeover Marker records paused automation, a Curb Lighthouse shows aggregate status without personal identity, and a Human-Machine Transfer Salon explains rail, walking, wheelchair, freight and staffed service [data:visual/assets/public-landmarks.json].

This is a public narrative and operating entrance, not a registered trademark, building, enterprise partnership or field result. If the mark cannot show an owner, a stop action and the ordinary route, the service window is withdrawn while the human path remains.

Railway memory—public recognition—AI return: one withdrawable spatial narrative

The cultural line uses no uncleared historical image, and it does not treat company exposure or model rank as “AI culture.” It joins three civic capabilities: Jing-Zhang railway continuity, engineering care and public passage as public memory; Zhongguancun reproducible methods, open collaboration and failure repair as the contribution rule; and the returned ordinary route, public receipt and next review date after every trial as AI public return. Public sources establish only the open-call and park context; interpretation, display and construction still require later rights and professional review [source:JINGZHANG-PUBLIC-FEEDBACK] [source:JINGZHANG-PARK-CATALOG].

Each landmark carries one withdrawable recognition receipt. The Zhongzhiyuan Curb Lighthouse shows version, gate, accountable role, pause reason, repair state and review date. The AI Origin Human-Takeover Marker records rejection, return to staffed service and completed repair. The Dazhongsi Transfer Receipt Wall places rail, walking, wheelchair, freight and staffed service in one complete journey. Every display must state the contribution, evidence receipt, independent-review role, rights/licence status, maintenance owner and expiry date. A missing field or any break in the ordinary route withdraws the display and reopens review. Five wayfinding layers run from site and public memory through ordinary route, status/consent, service/handover, and redress/archive/return. The yearly cycle repeats Q1 calibration, Q2 complete-route audit, Q3 bounded staffed demonstration and Q4 publication of repair, withdrawal or conditional continuation. Its international line is Build to test; publish to repair; return to public use. [data:visual/assets/culture-honor-return-system.json]





Beijing's 2025 rules for automated-vehicle road testing and demonstration applications place L3+ activities within a joint municipal working mechanism and approved roads/areas [source:BEIJING-AV-TEST-2025]. The 2025 policy-first-zone notice identifies designated roads, qualifications and time restrictions [source:BEIJING-AV-ROADS-2025]. The national MIIT/MPS/MOT framework requires an eligible test subject, driver, vehicle, testing evaluation and liability capacity [source:MIIT-AV-TEST-2021]. These sources establish a regulated pathway, not a local permit, road opening or operating baseline.

The package distinguishes known geometry-derived values, design_target trial gates, unknown local performance baselines and blocked conditions that prohibit expansion. The three research papers on street-canyon CFD, wind/heat/PM2.5 monitoring and roof/canyon geometry are method evidence only; they do not provide Jing-Zhang boundary conditions, health causality, accident rates or transferable percentages [source:LIU-URBAN-VENTILATION-2017] [source:MENG-WIND-HEAT-PM25-2022] [source:NOSEK-STREET-CANYON-2025].

1. Evidence hierarchy and autonomous-mobility decision boundary

The autonomy proposal first asks what action a source can support, then discusses vehicles, curbs or services. Registration, replay and design targets do not automatically become field facts.

Evidence level	Package examples	Can support	Cannot support
Task and regulatory pathway (formal/background)	Open-call, Beijing AV testing rules, designated-road notice and national framework	A regulated pathway, eligibility questions, authorization questions and professional review entry points	A project permit, an opened road, an operator or operating/incident performance
Provisional spatial base	Site/key areas, roads, autonomy nodes and test yards	Concept relationships, ordinary routes, node sequence and a whole-package recomputation trigger	Statutory redlines, road sections, speed/capacity, ownership, fire or accessibility compliance
Known and derived package readouts	GeoJSON, metrics.json, scenario cards, readiness and accountability matrices	File consistency, node actions, phase dependencies, issue lists and design gates	Existing footfall, takeover latency, conflict rate, vehicle performance or development intensity
Synthetic replay and design targets	AV-T01—T03, tabletop, known/design_target metrics	Stop, human takeover, rollback, data minimisation and next-test contracts	Field safety, public acceptance, staffing, P1/P2 authorization or official score
Papers and methods	Wind/heat/CFD, ecology, asset and embodied-intelligence methods	Future measurement variables, model inputs, uncertainty and professional review questions	Jing-Zhang wind field, health causality, accident rate, transferable percentage or permit basis
Blocked and accountability records	blocked conditions and insurance/safety/privacy/human-route gaps	When to stop, who must supply evidence and when scale-up is prohibited	Replacing accountability, approval, insurance or implementation results with "pending"

The review rule is: known means only that a value is readable from package files; design_target is not field performance; unknown must not be guessed; blocked stops scale-up. A local PASS proves only that a structure or state machine can be replayed; it is not a field, professional or implementation finding.

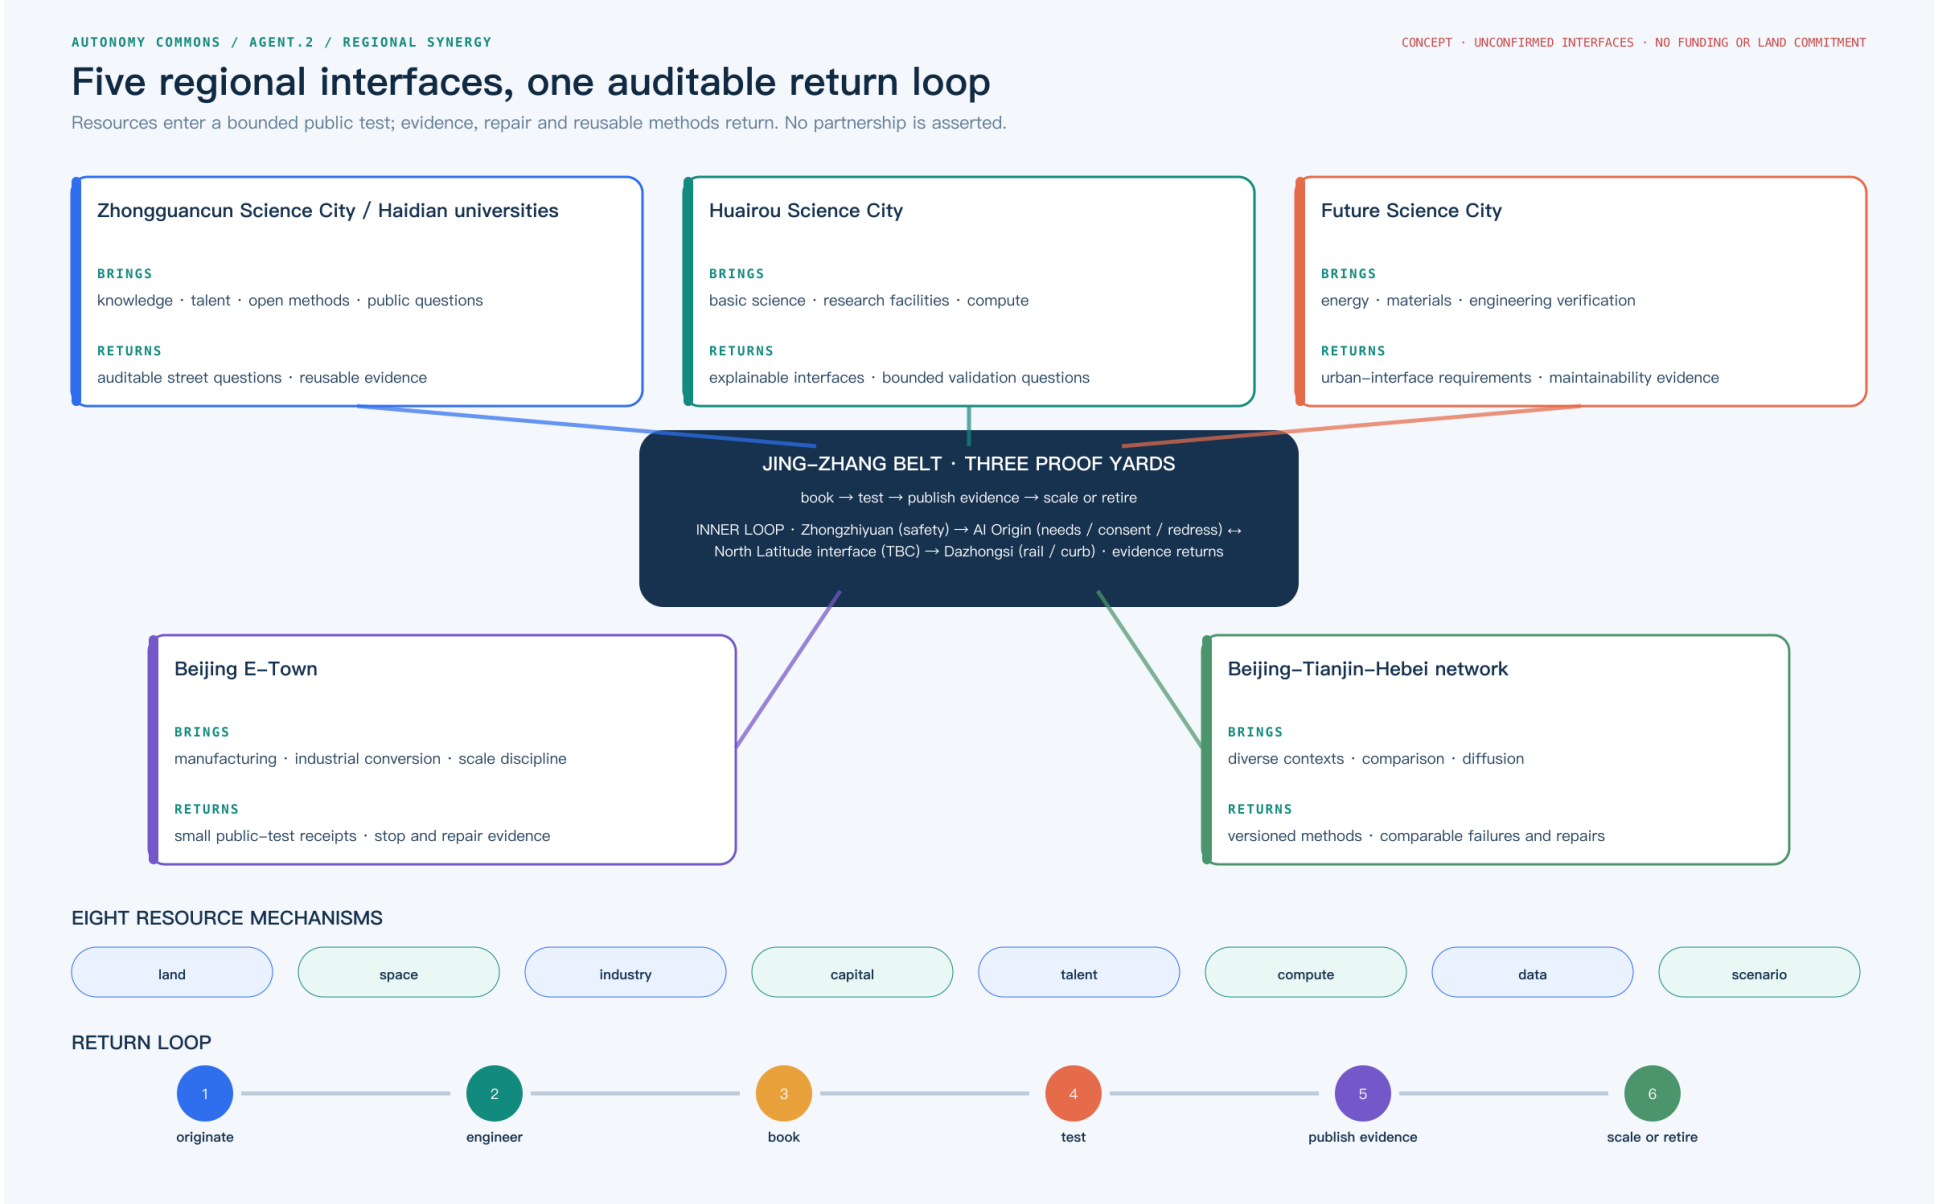
Coordinated Research Area: Industry and Future City Research

The coordinated research area connects AI industry, talent, rail, community, public governance and future-city research through public value rather than a single vendor. Enterprises may provide controlled test capacity, universities and students may participate in audit, communities provide real needs and redress, and rail and staffed service desks provide non-app entry points [source:AGENT-TASKBOOK]. The design intent is to make automation usable by residents, handover-ready for maintainers and reviewable by professionals. Spatially this assigns roles to the public axis and three yards; metrically it prioritises accessible continuity, equivalent human service, rollback and complaint closure, not vehicle count or model accuracy. Liability agreements, budgets, insurance, staffing and demand distribution remain unknown, so the industry chain is a coordination hypothesis rather than a partnership or implementation claim.

Five regional interfaces: transfer mechanisms, not partnership claims

Five unconfirmed regional interfaces each state what enters and what returns. Zhongguancun Science City and Haidian universities bring knowledge, talent, open methods and public questions. Huairou Science City brings basic science, research facilities and compute. Future Science City brings energy, materials and engineering verification. Beijing E-Town contributes manufacturing and scale discipline. The Beijing-Tianjin-Hebei network contributes diverse contexts and comparison. The Jing-Zhang Belt accepts only booked, bounded tests and public evidence, returning reusable methods, failure/repair records and maintainability requirements. The eight mechanisms—land, space, industry, capital, talent, compute, data and scenario—remain interfaces to confirm; they assert no partner, funding, land allocation or government commitment [data:visual/assets/regional-ecosystem.json].

An explicit public loop sits inside the belt. Zhongzhiyuan first turns incoming capability into bounded safety questions. The AI Origin Community adds resident needs, accessibility, consent, redress and equivalent human service to the test conditions. “North Latitude Community,” named by the taskbook, is shown separately as a community-coordination interface whose identity and extent await organiser confirmation; it records where questions go, why they are accepted or declined, and what remains unresolved. Dazhongsi then tests rail-curb-walking continuity. Evidence, repairs and unresolved items return to the five external interfaces along the same path. This proposal does not treat AI Origin and North Latitude as the same entity, nor infer resident representation, partnership or operating authority. Until the organiser confirms identity and accountability, North Latitude remains a taskbook needs interface only [data:visual/assets/regional-ecosystem.json#internal_public_return_loop].



Five regional interfaces, eight resource mechanisms and an auditable return loop

2. Spatial structure: one public axis, three test yards, two safety nets

The proposal is organised as one public axis + three test yards + two safety nets + twelve scenario cards. The public axis is the Jing-Zhang heritage park and its slow-mobility/blue-green connections. The test yards are Zhongzhiyuan, the AI Origin Community and Dazhongsi. The safety nets are:

- a human safety net: continuous accessible routes, staffed service, legible curbs, low-speed operation, emergency stop and human takeover;
- an ecology-and-data safety net: weather and network rollback, dark-sky and bird protection, data minimisation, public algorithm records and revocable consent.

Autonomous driving is a constrained service layer over walking, cycling, rail, transit, emergency and maintenance systems. Every vehicle or robot yields to the continuous human route. Curbs register who may stop, when, for how long and who clears the space [standard:BEIJING-ACCESSIBILITY-REGULATION] [standard:ISO-TR-4448-PUBLIC-MOBILE-ROBOTS].

Overall Design Area: Urban Renewal and Regulatory-Plan-Level Urban Design

The overall design area keeps provisional land-use, building, road, green, public-space and phasing layers as one spatial base; autonomous nodes express test relationships, not statutory redlines. Human routes, fire clearance, maintenance access, quiet interfaces and reversible upgrades come before service windows [depth:overall_spatial_structure]. Geometry is read back from land_use.geojson, buildings.geojson, constraints.geojson and phasing.geojson, while area, footprint, green and public-space ratios are recomputed from the metric ledger. Official regulatory plans, ownership, utilities, fire review and a demolition schedule are absent, so no floor-area, height, demolition or investment conclusion is made.

Three-Level Scope Framework

The three levels ask how policy and industry become public-space capacity, how curbs, stations, slow mobility and maintenance form a continuous system, and how each test yard becomes small, reversible and auditable. The following table translates that framework into three provisional key-area roles.

Key area	Public role	First reversible test	Hard boundary
Zhongzhiyuan	safety evaluation, simulation, vehicle-road-cloud interfaces and enterprise services	low-speed shuttle or inspection within a managed window	no social-road operation without approval; simulation is not safety proof
AI Origin Community	community service, accessible ride, public explanation and opt-out	assisted reservation with an equivalent human path	no app-only access, continuous resident tracking or forced consent
Dazhongsi	rail interchange, curb logistics and event-day human-machine separation	delivery/maintenance robot at a station-edge curb	no blocking fire, accessible or emergency routes; immediate rollback in rain, crowds or network loss

The three detailed points are stored in visual/assets/autonomy_nodes.json: AUTO-NODE-001 Zhongzhiyuan safety yard, AUTO-NODE-002 AI Origin accessible-service yard and AUTO-NODE-003 Dazhongsi human-machine transfer living room. They are provisional design markers, not designated test roads or statutory station locations [data:visual/assets/autonomy_nodes.json#AUTO-NODE-001].

Detailed Design of Key Areas

Zhongzhiyuan hosts managed-window safety evaluation, the AI Origin Community hosts equivalent automated and human service, and Dazhongsi hosts rail interchange, curb logistics and event-day separation. Each node must connect to a scenario card, curb state, responsible party, data boundary and stop condition; the node file is a low-confidence design register, not a permit or construction location. Survey, signal timing, station accessibility, fire routes, logistics windows and stratified resident baselines are still missing, so any first test must be closable, human-handover-ready and publicly replayable [data:visual/assets/autonomy_nodes.json#AUTO-NODE-001].

Public Interfaces and Reversible Relationships for the Three Test Yards

An autonomous node cannot be reduced to a vehicle stopping point. The ledger below translates each test yard's automation relationship into public interfaces, ordinary routes and reversible service relationships for later urban-design comparison; it does not change the node file or create a new redline. Official controls, tenure, utilities, fire review and a demolition schedule are absent, so no development control, storey count or engineering scale is filled in.



Key-area public roles and stop conditions

Caption: The key-area board is not a point list: it places enterprise service, equivalent community service and rail interchange at ground interfaces, and shows when the system returns to human, accessible and emergency routes.

4. Design rules for an automated future

Curb before vehicle. Draw the continuous human route and the exit-capable service zone before drawing a vehicle route. Curb states are open, booked, service, restricted and human-only; every change needs an owner, time window, sign and restoration condition.

The “status board plus removable functional bands” is a methodological transfer, not a local finding: curb research prompts fields for enforcement, public communication, data management and interagency coordination; citizen research on driverless streets treats greenery, seating, shelter and micromobility as questions for public preference in flexible zones; place-based eHMI research suggests that autonomous interfaces should be integrated with the context of a place. These studies only define what should be asked, measured and professionally reviewed next; their samples and outcomes are not transferred into Jing-Zhang preferences, dimensions or safety effects [source:DIEHL-CURBSPACE-MANAGEMENT-2021] [source:FAYYAZ-DRIVERLESS-STREETS-2025] [source:WANG-PLACE-BASED-EHMI-2025].

A typical curb cross-section: continuity before the service window

Cross-section band	Spatial move	Visible state and takeover	Fact that cannot be assumed yet
Continuous walking/accessibility band	Carry walking, wheelchair and care routes through station, service desk and crossings; vehicles do not enter	open or human-only; staff redirect around obstruction, water or blockage	Width, slope, tactile paving and lift status require a professional walk audit
Curb service band	Place short-term pick-up, loading, maintenance and clearing windows outside the public route	booked, service or restricted; a board shows owner, window and restoration condition	Right-of-way, ownership, fire clearance and conflict rate lack field baselines
Transfer/help node	Put station entrance, waiting, staffed help and complaint entry on one legible interface	Any service failure switches to human-only; the app is not the only entry	Station flow, waiting space and public experience require field verification
Maintenance/emergency fallback band	Keep removable elements, clearing actions and emergency access; do not turn a temporary service into a fixed road	Staff can emergency-stop, withdraw and restore public passage	Fire, utilities, maintenance shifts and emergency contacts require authorised confirmation

This cross-section fixes spatial relationships and state changes only. It does not fill engineering dimensions, road redlines, vehicle speeds or autonomous-operation permissions; those fields remain unknown pending survey, professional review and public participation.

Speed and rights are gates. The package does not invent a statutory speed limit. It records low-speed operation, yielding, stopping distance and human intervention as test fields [metric:autonomy_trial_speed_limit_status] [metric:remote_stop_response_seconds]. A successful trial lets the slowest pedestrian, wheelchair user, child and maintenance worker complete the route safely.

Someone can always take over. Every test yard has staffed service, physical emergency stop, remote stop, broadcast, incident log and withdrawal route. Uncertainty, communication loss, severe weather, crowds, ecological protection or blocked accessibility defaults to human service or shutdown [source:BEIJING-AV-TEST-2025] [source:MIIT-AV-TEST-2021].


Data is collected only to complete the service. Read the authorised curb state, obstacle class, accessible route and emergency message; do not build resident profiles or publish continuous camera streams. Public records show aggregate events, responsibility and corrections; retention and deletion require professional and legal confirmation [source:CASE-HELSINKI-AI-REGISTER] [source:CASE-UK-ATRS] [source:NIST-HUMAN-CENTERED-AI].

4.5 Public-route continuity: keep people moving before opening a service window

The first public question is not whether a vehicle completes a task, but whether an ordinary resident, wheelchair user, carer or maintainer still has a route that is visible, handover-ready and exit-capable. public-route-continuity.schema.json places three candidate nodes and four reviewer classes—professional, operator, resident and accessibility user—in one record. gap_ratio is a design demonstration; max_gap_ratio=0.25 is not a road-safety standard, and any blocking gap or invisible handover reopens the node [data:visual/assets/public-route-continuity.schema.json] [data:visual/assets/example-public-route-continuity.json].

The order is fixed: check all four reviewer readings at every node; check route continuity and visible human handover; only then allow an explainable caution state within the design limit. When it fails, the action is not to raise a model score: freeze new trials, keep the non-AI equivalent route open, publish the trigger and owner, repeat the affected node with every reviewer class, and require two consecutive clear rounds before reconsidering a window. accessible_route_continuity_ratio and autonomy_fallback_success_ratio remain unknown; the synthetic record proves field wiring and rejection paths only, not accessibility performance, resident outcomes, vehicle performance or permission [metric:accessible_route_continuity_ratio] [metric:autonomy_fallback_success_ratio].

The dependency-free checker and ten deterministic positive/negative fixtures are visual/assets/check_public_route_continuity.js and run_public_route_continuity.js. They make “the ordinary route remains” a public mechanism that a reviewer can deliberately break and observe failing; a local PASS is not field evidence.



All nine roles have an offline entrance, human help, revocable consent and a public complaint route; no one loses the service for not using an app [data:visual/assets/persona-role-matrix.json] [source:BEIJING-ACCESSIBILITY-REGULATION].

7. Landmarks and cultural narrative


The concept proposes three public landmarks: the Human-Takeover Marker at the Origin Community, recording rejected and paused automation; the Curb Lighthouse at Zhongzhiyuan, showing aggregate status and review gates without personal data; and the Human-Machine Transfer Salon at Dazhongsi, mapping rail, walking, wheelchair, freight and public service as one arrival system. These are cultural and interpretive proposals, not approved structures or sponsorship commitments [metric:ai_landmark_count] [data:visual/assets/public-landmarks.json].

Transport, Rail, Municipal Infrastructure, and Public Services

Walking, wheelchair, cycling, rail interchange, fire access, maintenance and public-service continuity take priority; an autonomous vehicle is a service layer and does not create road capacity. The first transport audit should locate the last 300–800 metre breaks, crossing conflicts, curb occupation, station accessibility and event-day spillover, then use counts and walk audits before any trial. Road ownership, signal timing, cross-sections, transit demand, loading, utilities and fire verification are still missing; the submitted roads and public-space layers therefore express relationships and candidate audit points, not engineering alignments or capacity [source:BEIJING-WALK-CYCLE-DB11-1761].

Land Use, Building Scale, and Retain-Renovate-Demolish Strategy

The package retains the provisional land-use and building base and does not convert autonomous nodes into development intensity, building-height, demolition or investment commitments. Reversible upgrades should use existing roads, stations, public buildings, shade and service desks; entrances, accessible toilets, waiting, maintenance/charging and human takeover are public-service interfaces. land_use.geojson, buildings.geojson, public_space.geojson and constraints.geojson express spatial relations, while area and footprint are recomputed from metrics.json. Ownership, structures, equipment, underground space, parking and cost data remain missing [data:geometry/land_use.geojson#LU-001].

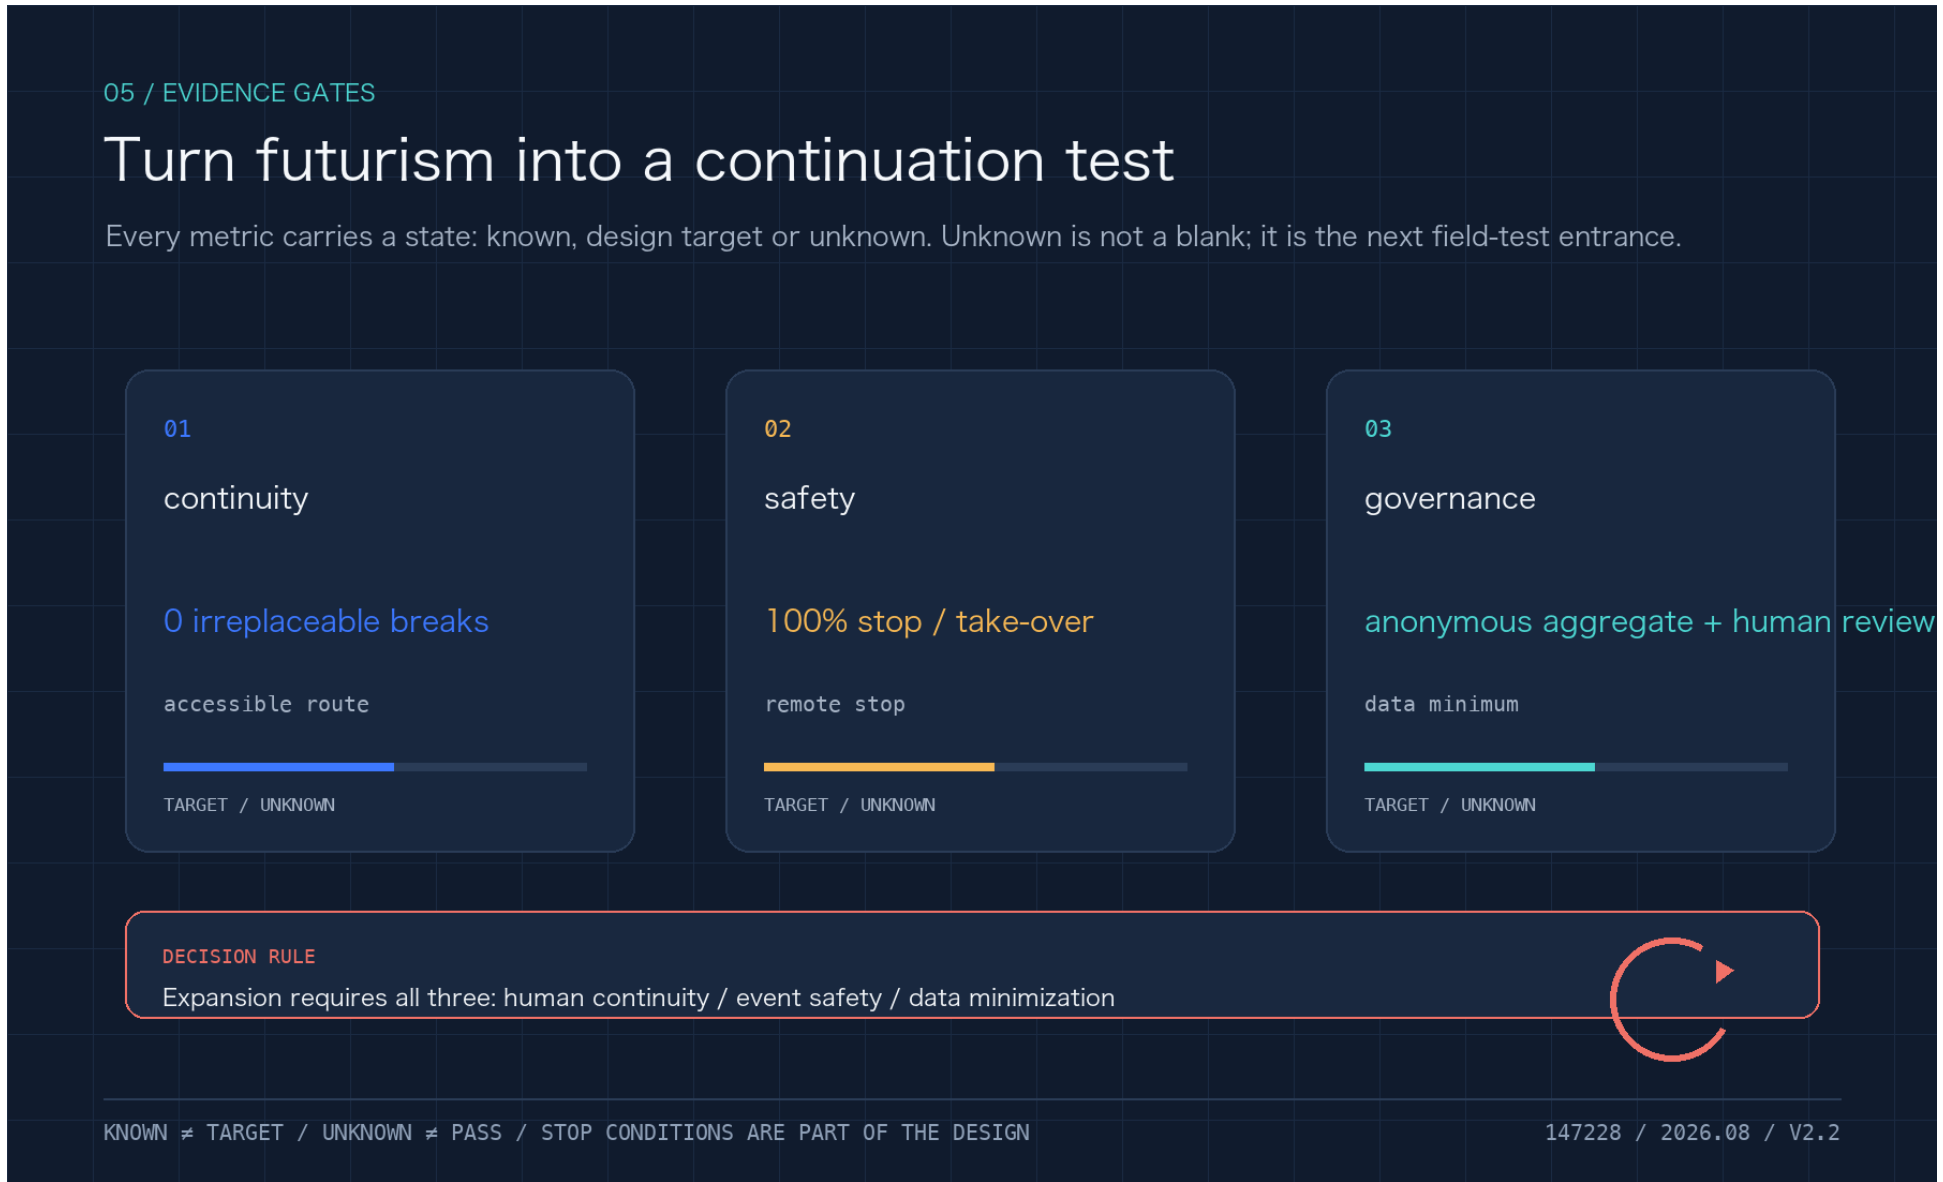


The item-level taskbook coverage, recomputed counts and evidence routes are in visual/assets/taskbook_coverage.json: 12 AV cards, 14 operation rows, 9 design roles, 3 public landmarks, 6 mechanism cases and 8 public components are each linked to a file. These counts prove deliverables exist; they do not prove event approval or an operating service.

8. Spatial layers, phases and evidence

The package retains the provisional base layers for land use, buildings, roads, green space and public space and adds visual/assets/autonomy_nodes.json. Existing area values are recomputed from the submitted layers [metric:site_area_sqm] [metric:green_ratio] [metric:public_space_ratio]. The new autonomy metrics are deliberately unknown or design_target: curb conflict rate, accessible-route continuity, remote-stop response, rollback success, data-minimisation coverage and trial-route length [metric:curb_conflict_rate] [metric:remote_stop_response_seconds] [metric:data_minimization_coverage_ratio].

Phasing is: P0 legible curbs and accessibility audit; P1 approved, low-speed, reversible tests; P2 conditional expansion only after safety, traffic, ecology, privacy, participation and liability gates pass. Automated vehicles are a service layer, never an assumed new road capacity. Climate, drainage, microclimate and ecology outcomes remain unknown without local observations and professional models [source:BEIJING-VENTILATION-NETWORK-2035] [source:BEIJING-FLOOD-PLAN-2021-2025] [source:BEIJING-BIRD-BIODIVERSITY-2024].

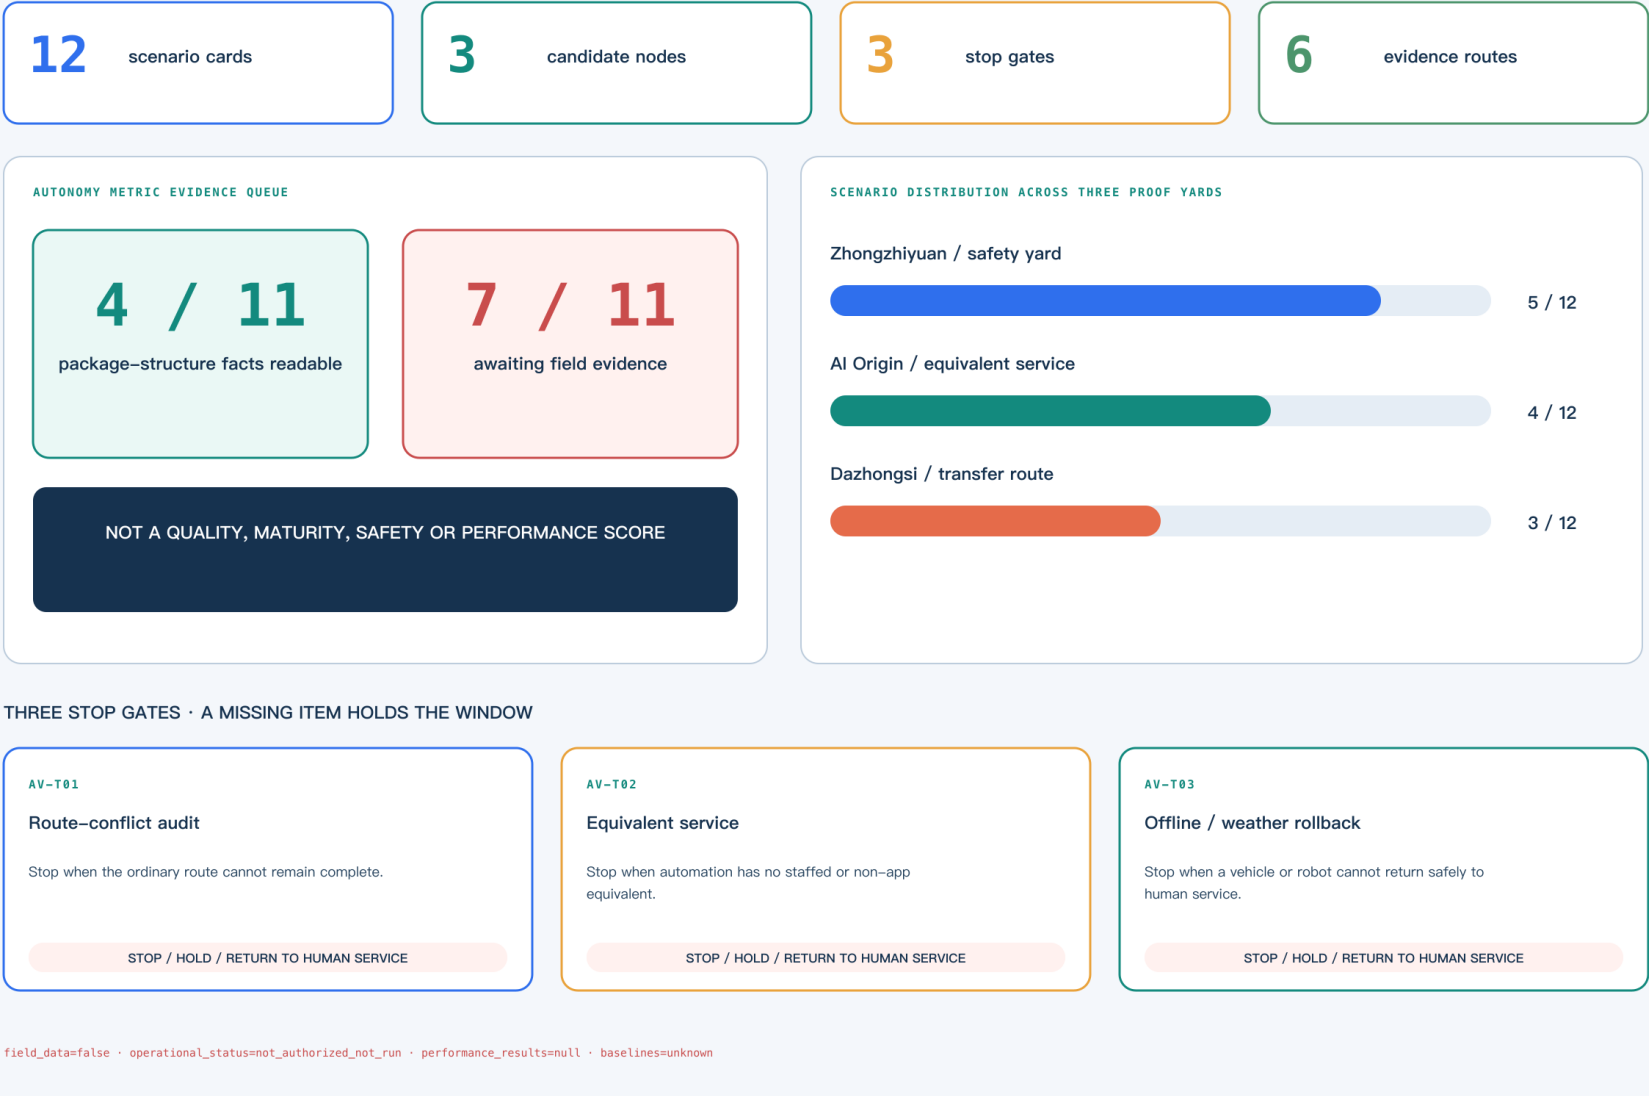


Stage gates, metrics and rollback

Caption: The metrics board separates the P0 readable-curb, P1 restricted-trial and P2 conditional-expansion gates, while showing unknown, design_target and rollback as different states.

Readiness is an evidence queue, not a score

Twelve scenario cards remain unauthorised and not run; separately, 7 of 11 metric items await field evidence.



Autonomy evidence queue: scenarios, nodes, gates, evidence routes and metric status

Caption: The evidence queue connects scenario cards, proof yards, gates and evidence routes. 7 / 11 means seven autonomy-specific metrics await field evidence; it is not a project-quality, maturity, safety, readiness or performance score.

Metrics, Area Recalculation, and Compliance Matrix

compliance_matrix.json covers announcement sections 1.3–1.5 and agent.1–agent.6. standard_matrix.json covers urban design, detailed planning, walking/cycling, accessibility, asset management, service-robot and automated-driving boundaries. design_depth_matrix.json links every depth item to narrative, layer, metric and rollback condition. No claim substitutes for a permit, vehicle certification, safety assessment, traffic review, insurance, privacy impact assessment, ecological review or construction document [source:BEIJING-AV-SAFETY-ASSESSMENT-2025].

Risk, Copyright, and Compliance

The package does not replace road-testing permission, vehicle certification, traffic organisation, fire review, insurance, privacy impact assessment, ecological review or construction documents. Rights and reuse boundaries remain in the package copyright statement and source records; future figures must be regenerated from the same structured data and must not turn a target into a known fact. Official polygons, road/utility/ownership, traffic, weather, drainage and ecology baselines, plus professional review, remain prerequisites for any operational or implementation conclusion.

The package-root risk.json records eight standard dimensions—policy and data uncertainty, spatial dispute, equity and inclusion, implementation complexity, technology maturity, public acceptance, privacy and operations cost—with bounded mitigations and human-review routes. It is a package-level boundary record, not a field risk assessment, permit or operational-safety conclusion.

Once official polygons, road/utility/ownership, traffic, weather, drainage and ecology baselines arrive, all layers, metrics, drawings, reports and self-check outputs must be regenerated. A single image must never be changed to turn a future target into a known fact.

References

The complete source index, publishers, use boundaries, access dates and limitations are recorded in sources.json; this section does not repeat the structured index, while the readable narrative keeps representative anchors next to the claims they support [source:SOURCE-REGISTRY]. The source layer exists to make each policy, standard, case and method traceable to its publisher, scope, licence boundary and known gap rather than treating citation count as quality. Geometry, areas, green/public ratios and autonomy metrics are read back from GeoJSON, JSON and matrices; papers define future measurement variables but do not provide local observations. Until official polygons, transport/ownership data, field baselines and professional review exist, trial, performance, permit and implementation outcomes remain targets, unknowns or pending formal evidence.

Boundary statement: this is an auditable concept and trial framework. It is not a government-approved plan, road-opening notice, autonomous-driving permit, company partnership, health/air-quality proof or construction commitment.